

Performance Analysis of Machine Learning Software

October 31, 2019

1 Performance analysis of Machine Learning (ML) software

```
[1]: #import modules
import pickle
import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import statsmodels as sm
import seaborn as sns
```

1.1 Acquire Data

We have 2 Dataframes: * Test info * Test results

```
[2]: #import dataframes
test_info=pd.read_csv("./TestInfo.csv")
file = open("./TestResults.pickle",'rb')
test_results=pickle.load(file,encoding='latin1')
```

The dataframe have in common one field= "TestId", we can join them

```
[3]: df=pd.merge(test_info, test_results)#.sort_values(by=["TestId"])
df.head()

fig=0
```

We can study what is inside the dataframe

```
[4]: df.describe()
```

```
[4]:
```

	TestId	CPUFrequency (MHz)	Threads	Build	Time (ms)	\
count	47.000000	47.000000	47.000000	47.000000	47.000000	
mean	23.000000	1297.872340	4.148936	7.127660	128.801418	
std	13.711309	462.267268	1.366982	3.187046	149.035930	
min	0.000000	1000.000000	1.000000	1.000000	19.000000	
25%	11.500000	1000.000000	3.500000	4.500000	51.500000	
50%	23.000000	1000.000000	5.000000	8.000000	83.333333	

75%	34.500000	2000.000000	5.000000	10.000000	114.500000
max	46.000000	2000.000000	5.000000	10.000000	800.000000

	PeakMemory (MB)
count	47.000000
mean	383.234043
std	125.454168
min	50.000000
25%	309.000000
50%	450.000000
75%	454.500000
max	460.000000

```
[5]: df.head()
```

```
[5]:
```

	TestId	Device	CPUFrequency (MHz)	Threads	MLNetwork	Build	Optimised	\
0	17	Device_0	1000	5	AlexNet	9	N	
1	16	Device_0	1000	5	AlexNet	8	N	
2	39	Device_1	1000	3	AlexNet	10	N	
3	31	Device_0	1000	5	AlexNet	10	N	
4	30	Device_0	1000	4	AlexNet	10	N	

	Time (ms)	PeakMemory (MB)
0	102.000000	449
1	104.000000	453
2	333.333333	302
3	100.000000	449
4	125.000000	450

```
[6]: print("Devices:",df["Device"].unique())
print("MLNetworks:",df["MLNetwork"].unique())
print("Optimization:",df["Optimised"].unique())
print("CPUFrequencies (MHz):",df["CPUFrequency (MHz)"].unique())
print()
print("Min and Max of Builds:",min(df["Build"]),"-",max(df["Build"]))
print("Min and Max of TestId: ",min(df["TestId"]),"-",max(df["TestId"]))
print("Min and Max of Peak Memory: ",min(df["PeakMemory_↵
↵(MB)"]),"-",max(df["PeakMemory (MB)"]))
print("Min and Max of Time: ",min(df["Time (ms)"]),"-",max(df["Time (ms)"]))
print("Min and Max of Threads:",min(df["Threads"]),"-",max(df["Threads"]))
```

```
Devices: ['Device_0' 'Device_1']
MLNetworks: ['AlexNet' 'MobileNet']
Optimization: ['N' 'Y']
CPUFrequencies (MHz): [1000 2000]
```

```
Min and Max of Builds: 1 - 10
```

Min and Max of TestId: 0 - 46
Min and Max of Peak Memory: 50 - 460
Min and Max of Time: 19.0 - 800.0
Min and Max of Threads: 1 - 5

1.2 Correlation

To help with Time and Memory we can create a new column, named Performance that is higher when both time and Memory are lower

$\text{Performance}(1/(Bs)) = (1/(\text{Memory } Time)) 1000 * 1000$


```
[7]: df["Performance"]=(1/(df["PeakMemory (MB)"]*df["Time (ms)"]))*1000 *1000

df=df.sort_values(by=["Performance"],ascending=False)
```

```
[8]: #Classification
df1=df
df1.loc[df["Device"] == "Device_0", 'Device_int'] = 0
df1.loc[df["Device"] == "Device_1", 'Device_int'] = 1
df1.loc[df["MLNetwork"] == "AlexNet", 'ML_int'] = 0
df1.loc[df["MLNetwork"] == "MobileNet", 'ML_int'] = 1
df1.loc[df["Optimised"] == "Y", 'Optimised_int'] = 1
df1.loc[df["Optimised"] == "N", 'Optimised_int'] = 0
df1=df1.drop(["TestId","MLNetwork","Device","Optimised"],1)
alex=df[df["MLNetwork"]=="AlexNet"]
mobile=df[df["MLNetwork"]=="MobileNet"]
```

```
[9]: f, ax = plt.subplots(figsize=(10, 8))
corr = df1.corr()
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.
    ↪diverging_palette(220, 10, as_cmap=True),
        square=True, ax=ax)
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb7e21d46d0>
```




```
[10]: sm = pd.plotting.scatter_matrix(df1, alpha=0.9, figsize=(12, 12),
    ↪diagonal='kde')

#Change label rotation
[s.xaxis.label.set_rotation(40) for s in sm.reshape(-1)]
[s.yaxis.label.set_rotation(0) for s in sm.reshape(-1)]

#May need to offset label when rotating to prevent overlap of figure
[s.get_yaxis().set_label_coords(-1.0,0.0) for s in sm.reshape(-1)]
[s.get_xaxis().set_label_coords(0.0,0.0) for s in sm.reshape(-1)]

#Hide all ticks
[s.set_xticks(()) for s in sm.reshape(-1)]
[s.set_yticks(()) for s in sm.reshape(-1)]

plt.show()
```



1.3 Findings:

- 47 Tests
- 2 MLNetworks (AlexNet, MobileNet)
- 2 Devices (Device_0, Device_1)
- 2 CPU Frequencies (1000, 2000)
- 2 Optimised Status (Yes or No)
- 5 Use of Threads (from 1 to 5)
- Peak Memory in MB (from 50 to 460)
- Time in ms (from 19 to 800)
- From heatmap and scatter matrix we can see that there is **correlation** between:

- **High**
- Threads and Memory -> More Threads use more memory
- Time and Device -> Device_0 uses more time
- ML and optimization -> Only MobileNet is tried optimized
- Performance and ML -> MobileNet is performing better
- Performance and Optimization -> Optimizing impact a lot on performance
- **Low**
- CPUFreq and Memory
- Build and Optimization
- Build and Device
- Build and Time
- Build and Performance
- As it is a Performance analysis of Machine Learning (ML) software the main point would be **maximization of Performance, minimizing Time and Memory**, so we should look on Time and Memory

Performance is maximized if: * More Thread are used * CPU Freq is Higher * Build are higher
 * Time and Memory are lower * Mobile Net is used * ML is optimized

1.3.1 Comparing the two ML Networks

```
[11]: for col in df.columns.drop(["Performance", 'Device_int', 'ML_int', '
      ↳ 'Optimised_int']):
      plt.
      ↳ scatter(df[df["MLNetwork"]=="AlexNet"][col], df[df["MLNetwork"]=="AlexNet"]["Performance"], a
      ↳ = 0.7)
      plt.
      ↳ scatter(df[df["MLNetwork"]=="MobileNet"][col], df[df["MLNetwork"]=="MobileNet"]["Performance
      ↳ = 0.7)
      fig+=1
      plt.title("Fig."+str(fig)+" Performance")
      plt.ylabel("Performance")
      plt.grid(b=True)
      plt.xlabel(col)
      plt.legend(["Performance AlexNet", "Performance MobileNet"])
      plt.show()
```

Fig.1 Performance

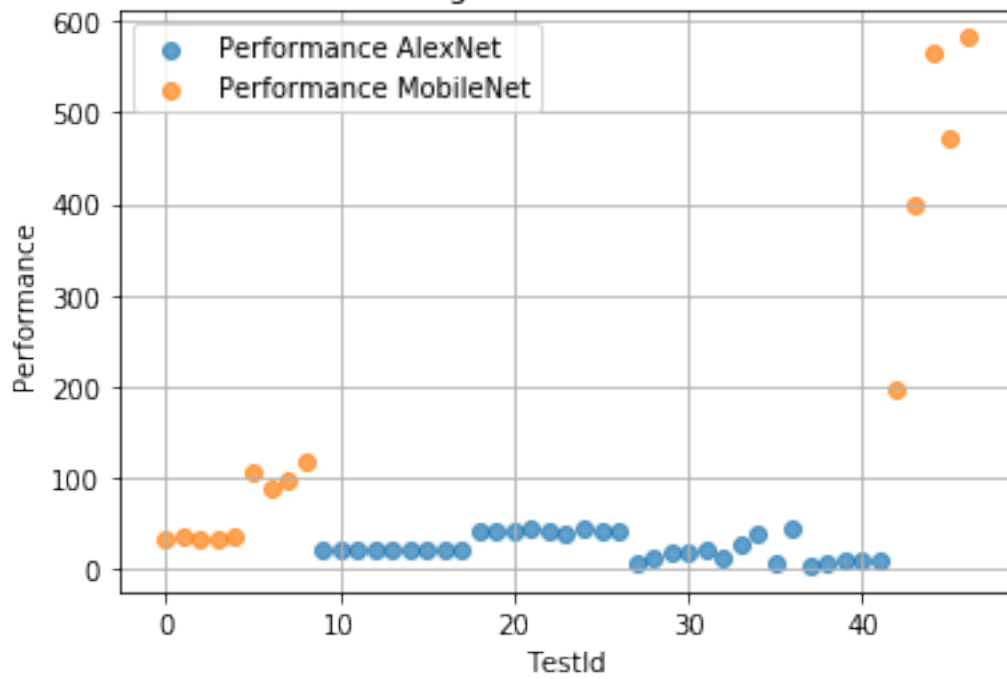


Fig.2 Performance




Fig.3 Performance

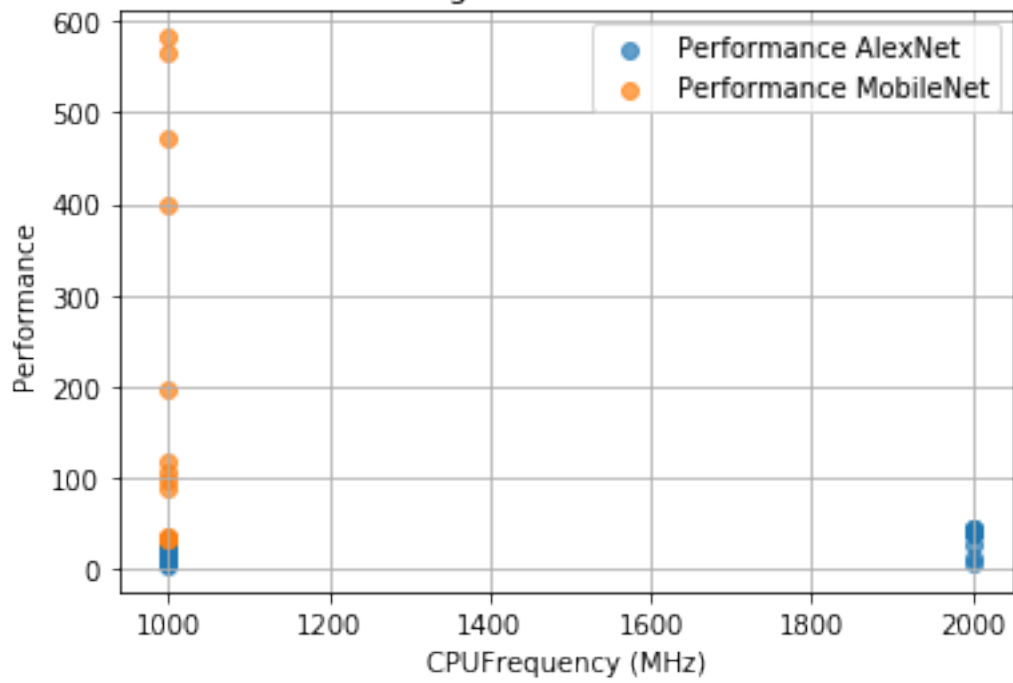
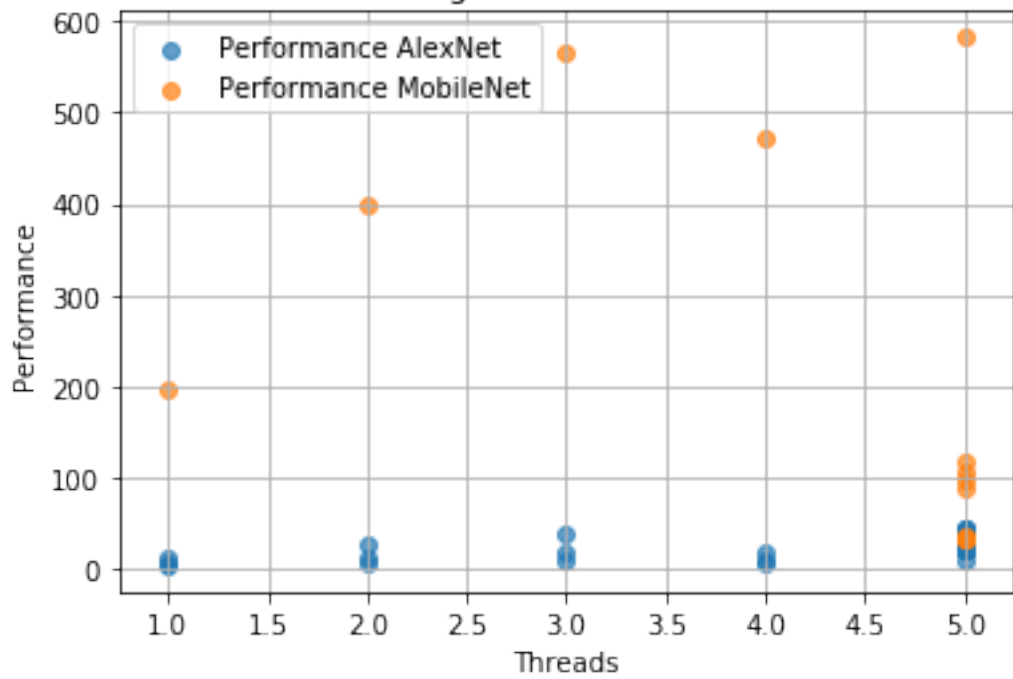


Fig.4 Performance



A scatter plot titled "Performance" comparing the performance of AlexNet and MobileNet. The y-axis is labeled "Performance" and ranges from 0 to 600. The x-axis has two categories: "AlexNet" and "MobileNet". The legend indicates that blue dots represent "Performance AlexNet" and orange dots represent "Performance MobileNet". AlexNet shows low performance values (0-50), while MobileNet shows a wide range of performance values (30-600).

Model	Performance
AlexNet	0
AlexNet	10
AlexNet	20
AlexNet	30
AlexNet	40
AlexNet	50
MobileNet	30
MobileNet	80
MobileNet	90
MobileNet	100
MobileNet	110
MobileNet	120
MobileNet	200
MobileNet	400
MobileNet	470
MobileNet	570
MobileNet	590

A scatter plot comparing the performance of AlexNet (blue dots) and MobileNet (orange dots) across 10 builds. The y-axis represents Performance (0 to 600) and the x-axis represents Build (1 to 10). AlexNet performance remains relatively stable, while MobileNet performance increases significantly, especially after build 6.

Build	Performance AlexNet	Performance MobileNet
1	25, 30, 35, 40	30, 35, 40, 45
2	20, 25, 30, 35	30, 35, 40, 45
3	20, 25, 30, 35	30, 35, 40, 45
4	20, 25, 30, 35	30, 35, 40, 45
5	20, 25, 30, 35	30, 35, 40, 45
6	20, 25, 30, 35	100, 110, 120, 130
7	20, 25, 30, 35	80, 90, 100, 110
8	20, 25, 30, 35	90, 100, 110, 120
9	20, 25, 30, 35	110, 120, 130, 140
10	20, 25, 30, 35	200, 210, 220, 230

Fig.7 Performance

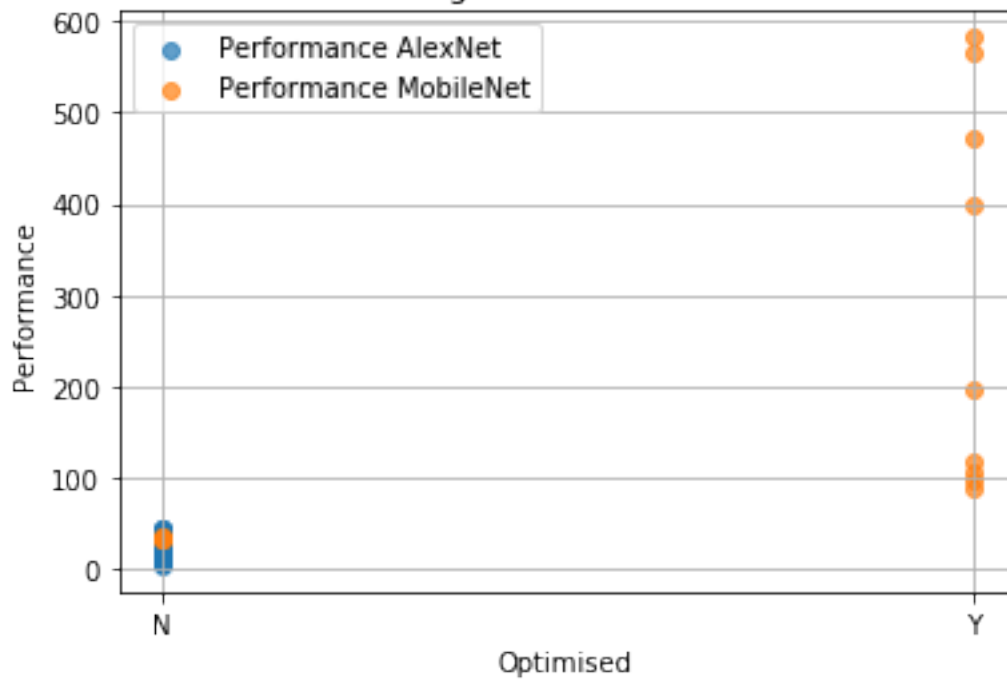

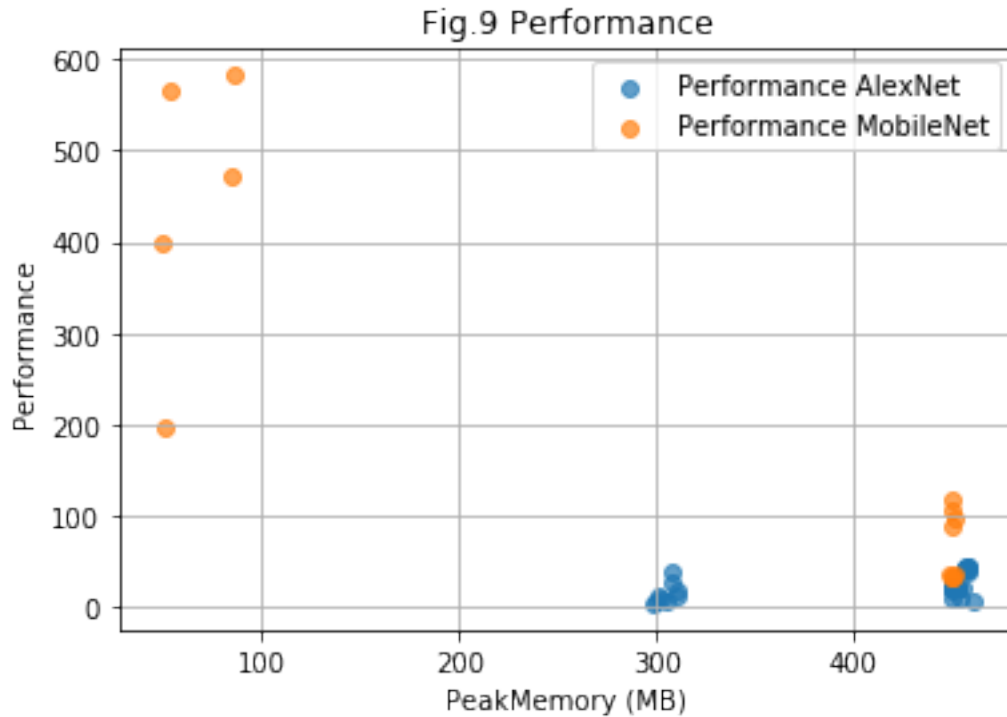


Fig.8 Performance





1.3.2 AlexNet

Findings

- AlexNet is not optimized
- More Threads uses more memory but less time
- More CPU Frequency means less time
- To minimize time and space we can check which value is closer to 0 in figure 13,14,15 Overall the best option performance wise is to use 5 threads, if memory is a constraint we can lower the number of threads

```
[12]: legend=[]
alex=alex.sort_values(by=["Threads"])

for o in alex["Optimised"].unique():
    for t in alex["Threads"].unique():
        legend.append(str(t)+" thread"+" Optimized: "+o)
for device in alex["Device"].unique():
    for cpu_freq in alex["CPUFrequency (MHz)"].unique():
        for optimised in alex["Optimised"].unique():
            for thread in alex["Threads"].unique():

                tmp=alex[alex["Threads"]==thread]
                tmp=tmp[tmp["CPUFrequency (MHz)"]==cpu_freq]
```

```

tmp=tmp[tmp["Optimised"]==optimised]
tmp=tmp[tmp["Device"]==device]
if not tmp.empty:
    plt.scatter(tmp["PeakMemory (MB)"],tmp["Time (ms)"])

if not tmp.empty:
    fig+=1;
    plt.title(("Fig."+str(fig)
              +" Device: "+device
              +" CPU Frequency (MHz): "+ str(cpu_freq)
              +" MLNetwork: AlexNet" ))
    plt.grid(b=True)
    plt.xlabel("PeakMemory (MB)")
    plt.ylabel("Time (ms)")
    plt.legend(legend)
    plt.show()

```

Fig.10 Device: Device_1 CPU Frequency (MHz): 1000 MLNetwork: AlexNet




Fig.11 Device: Device_0 CPU Frequency (MHz): 1000 MLNetwork: AlexNet





Fig.12 Device: Device_0 CPU Frequency (MHz): 2000 MLNetwork: AlexNet



1.3.3 MobileNet

Findings

- MobileNet is tested only on Device_0
- MobileNet is tested only on CPUFreq: 1000MHz
- Optimization is more efficient than not optimised
- Best case: 5 Thread, Optimised

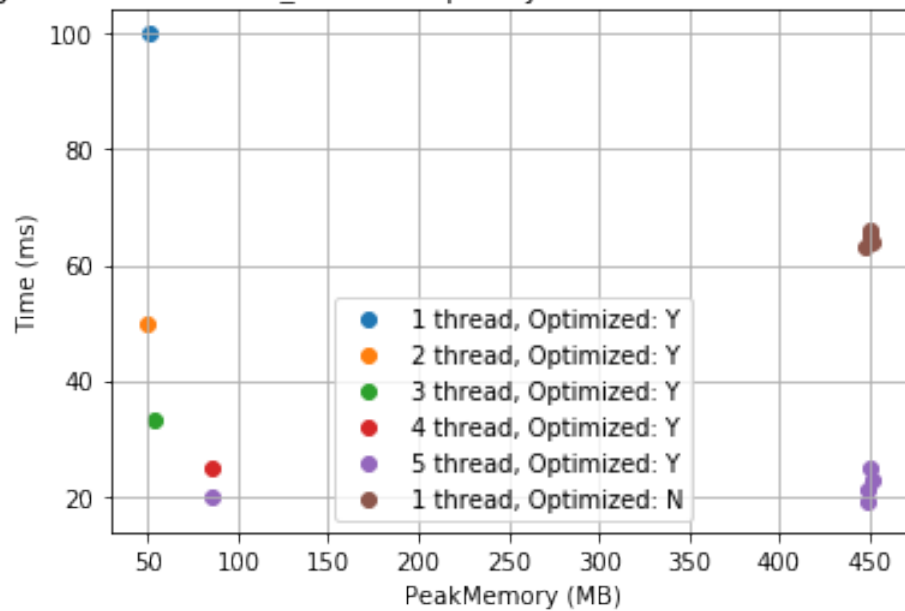
```
[13]: legend=[]
mobile=mobile.sort_values(by=["Threads"])
for o in mobile["Optimised"].unique():
    for t in mobile["Threads"].unique():
        legend.append(str(t)+" thread"+" Optimized: "+o)

for device in mobile["Device"].unique():
    for cpu_freq in mobile["CPUFrequency (MHz)"].unique():
        for optimised in mobile["Optimised"].unique():
            for thread in mobile["Threads"].unique():

                tmp=mobile[mobile["Threads"]==thread]
                tmp=tmp[tmp["CPUFrequency (MHz)"]==cpu_freq]
                tmp=tmp[tmp["Optimised"]==optimised]
                tmp=tmp[tmp["Device"]==device]
                if not tmp.empty:
                    plt.scatter(tmp["PeakMemory (MB)"],tmp["Time (ms)"])

if not tmp.empty:
    fig+=1;
    plt.grid(b=True)
    plt.title(("Fig."+str(fig)
              +" Device: "+device
              +" CPU Frequency (MHz): "+ str(cpu_freq)
              +" MLNetwork: MobileNet" ))
    plt.xlabel("PeakMemory (MB)")
    plt.ylabel("Time (ms)")
    plt.legend(legend)
    plt.show()
```

Fig.13 Device: Device_0 CPU Frequency (MHz): 1000 MLNetwork: MobileNet



1.4 Follow Up Actions

Talk with the engineers and show that * Tests are promising, they should try to test: * CPU Freq: 1000 and 2000MHz on Device_1 with MobileNet and 5 thread * AlexNet optimized

- MobileNet seems very promising as performance compared to AlexNet both for time and memory
- Optimization is fundamental, AlexNet should implement it
- Highest build generally perform better